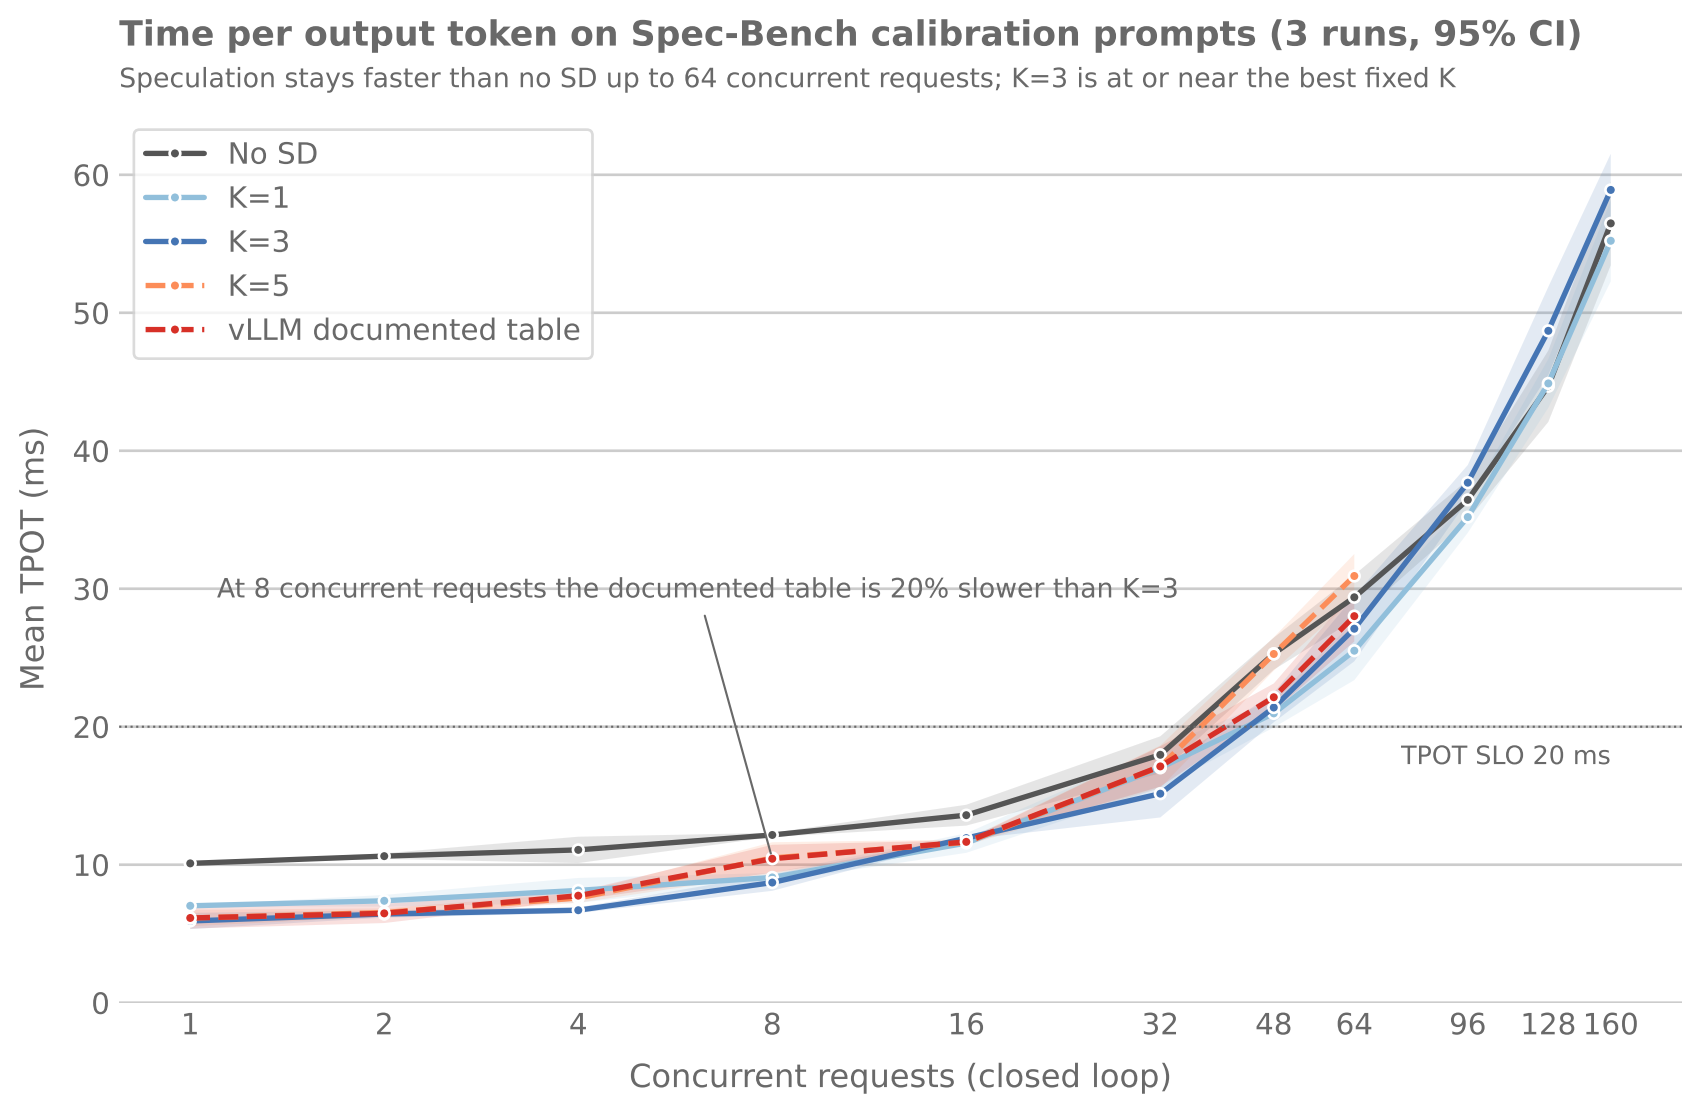
Time per output token on Spec-Bench calibration prompts (3 runs, 95% CI)
Speculation stays faster than no SD up to 64 concurrent requests; K=3 is at or near the best fixed K
No SD
K=1
K=3
K=5
vLLM documented table
Mean TPOT (ms)
0
10
20
30
40
50
60
At 8 concurrent requests the documented table is 20% slower than K=3
TPOT SLO 20 ms
1
2
4
8
16
32
48
64
96
128
160
Concurrent requests (closed loop)